

CMPE-283 Assignment - 1

Overview

In this assignment we will create a linux kernel module to query different MSR's to figure out different virtualization features available to the CPU.

For this assignment we have enabled a GCP VM based on linux operating system. We verify the MSR's by inserting a new module to the kernel and discover its features by the features it discovers.

Contribution

Asish's Contribution

Dheeraj and I collaborated over meeting rooms at SJSU. I began by setting up an virtual machine on GCP Cloud. After that, I set up my github on the vm to make updation and version control easy. Then i moved on with cloning the linux github repo the Makefile and the.c file to my virtual machine (VM). I then searched for various capability regions in SDM, as well as report_capability and the remaining msrs for each capability.

Dheeraj's Contribution

In a meeting, Asish and I went over the specifics of the task and watched the assignemtn 1 video that covered the requisites. I loaded the makefile and starter.c file into the VM provided by Asish, then I started adding the last four struct definitions for the various capability info areas. I added the four last msrs readings for each capability to the detect_vmx_features function, and then I added the calls to report_capability. I then attempted to run make and sudo insmod on the freshly created.ko file in order to do testing.

Procedure

1. Create a GCP account and avail your \$300 free credits with the sjsu.edu account.
2. After the \$300 credit is offered we activate the cloud shell and run the follwoing command

```
gcloud compute instances create cmpe283-vm4 --enable-nested-virtualization --zone=us-west4-b --machine-type=n2-standard-8 --network-interface=network-tier=PREMIUM,subnet=default --create-disk=auto-delete=yes,boot=yes,device-name=instance-1,image=projects/ubuntu-os-cloud/global/images/ubuntu-2004-focal-v20220204,mode=rw,size=200 --metadata=ssh-keys=asish:"ssh-ed25519 xxxxxxxxxxxx"
```

the following command will creat a GCP CM instance with 16 gb of ram and 200GB of disk space.

```
asishraju_vachavaya@cloudshell:~ (cmpe-283-407516) $ gcloud compute instances create cmpe283-vm2 --enable-nested-virtualization --zone=us-west1-b --machine-type=n2-standard-2 --network-interface=network-tier=PREMIUM,subnet=default --create-disk=auto-delete=yes,boot=yes,device-name=instance-1,image=projects/ubuntu-os-cloud/global/images/ubuntu-2004-focal-v20220204,mode=rw,size=10 --metadata=ssh-keys=asish:"ssh-ed25519 AAAAC3NzaC1lZDI1NTE5AAAAIK1ES7Oqifn41kxTP24F9bL6+9b0dy8ap5F"
FDZ/TIMB admin@USCS-Mac109
Created [https://www.googleapis.com/compute/v1/projects/cmpe-283-407516/zones/us-west1-b/instances/cmpe283-vm2].
WARNING: Some requests generated warnings:
- The resource 'projects/ubuntu-os-cloud/global/images/ubuntu-2004-focal-v20220204' is deprecated. A suggested replacement is 'projects/ubuntu-os-cloud/global/images/ubuntu-2004-focal-v20231130'.
NAME: cmpe283-vm2
ZONE: us-west1-b
MACHINE_TYPE: n2-standard-2
PREEMPTIBLE:
INTERNAL_IP: 10.138.0.2
EXTERNAL_IP: 35.197.70.119
STATUS: RUNNING
asishraju_vachavaya@cloudshell:~ (cmpe-283-407516) $
```

3. Now we ssh into our newly created instance and install git and other necessary tools using the command below

```
sudo apt-get install git
sudo apt-get install gcc
```

```
sudo apt-get install make
```

4. Once the following is done now we clone the linux repo from github

```
asish@cmpe283-vm2:~$ git clone https://github.com/dheerajnandigama/linux.git
Cloning into 'linux'...
remote: Enumerating objects: 9861292, done.
remote: Total 9861292 (delta 0), reused 0 (delta 0), pack-reused 9861292
Receiving objects: 100% (9861292/9861292), 4.61 GiB | 27.42 MiB/s, done.
Resolving deltas: 100% (8054415/8054415), done.
Updating files: 100% (82432/82432), done.
asish@cmpe283-vm2:~$ ;s
-bash: syntax error near unexpected token `;'
asish@cmpe283-vm2:~$ ls
linux
asish@cmpe283-vm2:~$ ls
linux
asish@cmpe283-vm2:~$ cd linux/
asish@cmpe283-vm2:~/linux$ ls
COPYING      Kbuild      MAINTAINERS  arch  crypto  include  ipc  mm  samples  sound  virt
CREDITS      Kconfig     Makefile    block  drivers  init      kernel  net  scripts  tools
Documentation LICENSES    README     certs  fs       io_uring  lib  rust  security  usr
asish@cmpe283-vm2:~/linux$
```

5. Then we clone our github repo

```
asish@cmpe283-vm2:~/linux$ git remote -v
origin  https://github.com/dheerajnandigama/linux.git (fetch)
origin  https://github.com/dheerajnandigama/linux.git (push)
asish@cmpe283-vm2:~/linux$
```

6. Now we check for available virtualization flags available in our operating system.

```
asish@cmpe283-vm2:~$ cat /proc/cpuinfo |grep vmx
flags       : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2
              ss ht syscall nx pdpe1gb rdtscp lm constant_tsc rep_good nopl xtopology nonstop_tsc cpuid tsc_known_freq pni pclmu
              lqadq vmx ssse3 fma cx16 pcid sse4_1 sse4_2 x2apic movbe popcnt aes xsave avx f16c rdrand hypervisor lahf_lm abm 3dno
              wprefetch invpcid_single ssbd ibrs ibpb stibp ibrs_enhanced tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase
              tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm avx512f avx512dq rdseed adx smap clflushopt clwb avx512cd avx51
              2bw avx512vl xsaveopt xsavec xgetbv1 xsaves arat avx512_vnni md_clear arch_capabilities
vmx flags    : vnmi preemption_timer invvpid ept_x_only ept_ad flexpriority tsc_offset vtptr mtf vapid ept vpid u
              nrestricted_guest vapid_reg shadow_vmcs
flags       : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2
              ss ht syscall nx pdpe1gb rdtscp lm constant_tsc rep_good nopl xtopology nonstop_tsc cpuid tsc_known_freq pni pclmu
              lqadq vmx ssse3 fma cx16 pcid sse4_1 sse4_2 x2apic movbe popcnt aes xsave avx f16c rdrand hypervisor lahf_lm abm 3dno
              wprefetch invpcid_single ssbd ibrs ibpb stibp ibrs_enhanced tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase
              tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm avx512f avx512dq rdseed adx smap clflushopt clwb avx512cd avx51
              2bw avx512vl xsaveopt xsavec xgetbv1 xsaves arat avx512_vnni md_clear arch_capabilities
vmx flags    : vnmi preemption_timer invvpid ept_x_only ept_ad flexpriority tsc_offset vtptr mtf vapid ept vpid u
              nrestricted_guest vapid_reg shadow_vmcs
```

7. Now we modify the `cmpe283-1.c` to add other MSR directories as mentioned below

```
struct capability_info pinbased[5] =
{
    { 0, "External Interrupt Exiting" },
    { 3, "NMI Exiting" },
    { 5, "Virtual NMIs" },
    { 6, "Activate VMX Preemption Timer" },
}
```

```
    { 7, "Process Posted Interrupts" }  
};
```

```
struct capability_info procbased[21] =  
{  
    { 2, "Interrupt-window exiting" },  
    { 3, "Use TSC offsetting " },  
    { 7, "HLT exiting " },  
    { 9, "INVLPG exiting " },  
    { 10, "MWAIT exiting" },  
    { 11, "RDPMC exiting" },  
    { 12, "RDTSC exiting" },  
    { 15, "CR3-load exiting" },  
    { 16, "CR3-store exiting" },  
    { 19, "CR8-load exiting" },  
    { 20, "CR8-store exiting" },  
    { 21, "Use TPR shadow " },  
    { 22, "NMI-window exiting" },  
    { 23, "MOV-DR exiting" },  
    { 24, "Unconditional I/O exiting" },  
    { 25, "Use I/O bitmaps " },  
    { 27, "Monitor trap flag " },  
    { 28, "Use MSR bitmaps" },  
    { 29, "MONITOR exiting" },  
    { 30, "PAUSE exiting" },  
    { 31, "Activate secondary controls" }  
};
```

```
struct capability_info secondary_procbased[23] =  
{  
    { 0, "Virtualize APIC accesses" },  
    { 1, "Enable EPT " },  
    { 2, "Descriptor-table exiting " },  
    { 3, "Enable RDTSCP " },  
    { 4, "Virtualize x2APIC mode" },  
    { 5, "Enable VPID" },  
    { 6, "WBINVD exiting" },  
    { 7, "Unrestricted guest" },  
    { 8, "APIC-register virtualization" },  
    { 9, "Virtual-interrupt delivery" },  
    { 10, "PAUSE-loop exiting" },  
    { 11, "RDRAND exiting " },  
    { 12, "Enable INVPCID" },  
    { 13, "Enable VM functions" },  
    { 14, "VMCS shadowing" },  
    { 15, "Enable ENCLS exiting " },  
    { 16, "RDSEED exiting " },  
    { 17, "Enable PML" },  
    { 18, "EPT-violation #VE" },  
    { 19, "Conceal VMX non-root operation from Intel PT" },  
};
```

```

    { 20, "Enable XSAVES/XRSTORS" },
    { 22, "Mode-based execution control for EPT" },
    { 25, "Use TSC scaling" }
};

```

```

struct capability_info entry_controls[9] =
{
    { 2, "Load debug controls" },
    { 9, "IA-32e mode guest" },
    { 10, "Entry to SMM" },
    { 11, "Deactivate dual-monitor treatment " },
    { 13, "Load IA32_PERF_GLOBAL_CTRL" },
    { 14, "Load IA32_PAT" },
    { 15, "Load IA32_EFER" },
    { 16, "Load IA32_BNDCFGS" },
    { 17, "Conceal VM entries from intel PT" }
};

```

```

struct capability_info exit_controls[11] =
{
    { 2, "Save debug controls" },
    { 9, "Host address-space size" },
    { 12, "Load IA32_PERF_GLOBAL_CTRL" },
    { 15, "Acknowledge interrupt on exit " },
    { 18, "Save IA32_PAT" },
    { 19, "Load IA32_PAT" },
    { 20, "Save IA32_EFER" },
    { 21, "Load IA32_EFER" },
    { 22, "Save VMX-preemption timer value" },
    { 23, "Clear IA32_BNDCFGS" },
    { 24, "Conceal VM exits from Intel PT" }
};

```

inside the `detect_vmx_features()` function we make these required changes to output the MSR capabilities.

```

/* Pinbased controls */
rdmsr(IA32_VMX_PINBASED_CTL, lo, hi);
pr_info("Pinbased Controls MSR: 0x%llx\n",
    (uint64_t)(lo | (uint64_t)hi << 32));
report_capability(pinbased, 5, lo, hi);

/* Procbased controls */
rdmsr(IA32_VMX_PROCBASED_CTL, lo, hi);
pr_info("Procbased Controls MSR: 0x%llx\n",
    (uint64_t)(lo | (uint64_t)hi << 32));
report_capability(procbased, 21, lo, hi);

```

```

    /* Secondary Procbased2 controls */
    rdmsr(IA32_VMX_PROCBASED_CTL2, lo, hi);
    pr_info("Secondary Procbased Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
    report_capability(secondary_procbased, 23, lo, hi);

    /* Entry controls */
    rdmsr(IA32_VMX_ENTRY_CTL2, lo, hi);
    pr_info("Entry Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
    report_capability(entry_controls, 9, lo, hi);

    /* Exit controls */
    rdmsr(IA32_VMX_EXIT_CTL2, lo, hi);
    pr_info("Exit Controls MSR: 0x%llx\n",
        (uint64_t)(lo | (uint64_t)hi << 32));
    report_capability(exit_controls, 11, lo, hi);

```

8. After the following modification is made we compile the C code using the make file provided.

```

sudo make
sudo cp cmpe283-1.ko
sudo dmesg

```

9. The Pin-Based VM-Execution Controls output

```

[15060.927917] Pinbased Controls MSR: 0x7f00000016
[15060.927918] External Interrupt Exiting: Can set:Yes, Can clear:Yes
[15060.927919] NMI Exiting: Can set:Yes, Can clear:Yes
[15060.927920] Virtual NMIs: Can set:Yes, Can clear:Yes
[15060.927921] Activate VMX Preemption Timer: Can set:Yes, Can clear:Yes
[15060.927922] Process Posted Interrupts: Can set:No, Can clear:Yes

```

10. The Processor based VM Execution Controls output

```

[15060.927924] Procbased Controls MSR: 0xffff9ffe0401e172
[15060.927925]   Interrupt-window exiting: Can set:Yes, Can clear:Yes
[15060.927926]   Use TSC offsetting : Can set:Yes, Can clear:Yes
[15060.927931]   HLT exiting : Can set:Yes, Can clear:Yes
[15060.927932]   INVLPG exiting : Can set:Yes, Can clear:Yes
[15060.927933]   MWAIT exiting: Can set:Yes, Can clear:Yes
[15060.927933]   RDPMC exiting: Can set:Yes, Can clear:Yes
[15060.927934]   RDTSC exiting: Can set:Yes, Can clear:Yes
[15060.927935]   CR3-load exiting: Can set:Yes, Can clear:No
[15060.927935]   CR3-store exiting: Can set:Yes, Can clear:No
[15060.927936]   CR8-load exiting: Can set:Yes, Can clear:Yes
[15060.927937]   CR8-store exiting: Can set:Yes, Can clear:Yes
[15060.927938]   Use TPR shadow : Can set:Yes, Can clear:Yes
[15060.927938]   NMI-window exiting: Can set:Yes, Can clear:Yes
[15060.927939]   MOV-DR exiting: Can set:Yes, Can clear:Yes
[15060.927940]   Unconditional I/O exiting: Can set:Yes, Can clear:Yes
[15060.927940]   Use I/O bitmaps : Can set:Yes, Can clear:Yes
[15060.927941]   Monitor trap flag : Can set:Yes, Can clear:Yes
[15060.927942]   Use MSR bitmaps: Can set:Yes, Can clear:Yes
[15060.927942]   MONITOR exiting: Can set:Yes, Can clear:Yes
[15060.927943]   PAUSE exiting: Can set:Yes, Can clear:Yes
[15060.927944]   Activate secondary controls: Can set:Yes, Can clear:Yes

```

11. The Entry control VM-Execution Control output

```

[15060.927963] Entry Controls MSR: 0xd3ff000011ff
[15060.927964]   Load debug controls: Can set:Yes, Can clear:No
[15060.927965]   IA-32e mode guest: Can set:Yes, Can clear:Yes
[15060.927966]   Entry to SMM: Can set:No, Can clear:Yes
[15060.927966]   Deactivate dual-monitor treatment : Can set:No, Can clear:Yes
[15060.927967]   Load IA32_PERF_GLOBAL_CTRL: Can set:No, Can clear:Yes
[15060.927967]   Load IA32_PAT: Can set:Yes, Can clear:Yes
[15060.927968]   Load IA32_EFER: Can set:Yes, Can clear:Yes
[15060.927969]   Load IA32_BNDCFGS: Can set:No, Can clear:Yes
[15060.927969]   Conceal VM entries from intel PT: Can set:No, Can clear:Yes

```

12. The Exit control VM-Execution Control output

```

[15060.927971] Exit Controls MSR: 0x7fefff00036dff
[15060.927972]   Save debug controls: Can set:Yes, Can clear:No
[15060.927972]   Host address-space size: Can set:Yes, Can clear:Yes
[15060.927973]   Load IA32_PERF_GLOB AL_CTRL: Can set:No, Can clear:Yes
[15060.927974]   Acknowledge interrupt on exit : Can set:Yes, Can clear:Yes
[15060.927974]   Save IA32_PAT: Can set:Yes, Can clear:Yes
[15060.927975]   Load IA32_PAT: Can set:Yes, Can clear:Yes
[15060.927975]   Save IA32_EFER: Can set:Yes, Can clear:Yes
[15060.927976]   Load IA32_EFER: Can set:Yes, Can clear:Yes
[15060.927977]   Save VMX-preemption timer value: Can set:Yes, Can clear:Yes
[15060.927977]   Clear IA32_BNDCFGS: Can set:No, Can clear:Yes
[15060.927978]   Conceal VM exits from Intel PT: Can set:No, Can clear:Yes

```