

CMPE-283 Assignment - 2

Overview

In this assignment we will modify the CPUID emulation code in KVM, for this assignment we will go with 2 leaf node that is `0x4FFFFFFD` and `0x4FFFFFFC` with some predefined rules for the output. For leaf nodes `0x4FFFFFFD` and `0x4FFFFFFC`, if `%ecx` (on input) contains a value not defined by the SDM, return 0 in all `%eax`, `%ebx`, `%ecx` registers and return `0xFFFFFFFF` in `%edx`. For exit types not enabled in KVM, return 0s in all four registers.

Contribution

Asish's Contribution

I created a new GCP-VM instance as we encountered some space issue due to building the kernel for linux, I cloned the torvalds/linux repository that Dheeraj had forked, I started by installing all the required tools to compile and build a kernel. Next, I modified `kvm_emulate_cpuid` to include the leaf function, and in the process, I read up on atomic variables. In order to test my code, I also installed Virt-Manager which would enable run a virtual manager inside VM.

Dheeraj's Contribution

I went over how to obtain and build the Linux kernel and spoke about the next steps needed to finish the assignment. I locally cloned and forked the Linux kernel repository in my development environment, then I started tinkering with the `cpuid.c` file in KVM by adding a leaf function to `kvm_emulate_cpuid`. I practiced creating the kernel and getting a nested virtual machine ready for testing in parallel. I configured GCP to provide display output as we would need it to run an virtualized ubuntu instance within our VM connected via chrom remote desktop.

Procedure

1. Created new VM in GCP cloud.

```
gcloud compute instances create cmpe283-vm2 --enable-nested-virtualization --zone=us-west4-b --machine-type=n2-standard-8 --network-interface=network-tier=PREMIUM,subnet=default --create-disk=auto-delete=yes,boot=yes,device-name=instance-1,image=projects/ubuntu-os-cloud/global/images/ubuntu-2004-focal-v20220204,mode=rw,size=200 --metadata=ssh-keys=asish:"ssh-ed25519 xxxxxxxxxxxx"
```

the following command will create a GCP CM instance with 16 gb of ram and 200GB of disk space.

2. Now we ssh into our newly created instance and install git and other necessary tools using the command below

```
sudo apt-get install git
sudo apt-get install gcc
sudo apt-get install make
```

3. Once the following is done now we clone the linux repo from github
4. Then we clone our github repo
5. We check the current installed kernel in our VM using the below command

```
uname -a
```

6. Once the required files are present we move our cwd to `cd linux` and run the following command

```
cp /boot/config-5.11.0-1029-gcp ~/linux/.config
```

7. `sudo nano .config` to edit the following lines

```
CONFIG_SYSTEM_REVOCATION_KEYS=""  
  
CONFIG_SYSTEM_TRUSTED_KEYS=""
```

8. now we will make our kernel using the config we just added, will press enter to default all values as this step will ask for many questions.

```
make oldconfig
```

9. once config has been built now we build our kernel modules and install the built kernel

```
make -j 8 modules && make -j 8
```

10. After the modules and kernel is built we install them

```
make -j 8 modules_install && make -j 8 install
```

```
asish@cmpe283-vm2:~/linux$ make -j2 && make modules -j2 && make install -j2 && make modules_install -j2  
mkdir -p /home/asish/linux/tools/objtool && make O=/home/asish/linux subdir=tools/objtool --no-print-directory -C o  
bjtool  
mkdir -p /home/asish/linux/tools/bpf/resolve_btfids && make O=/home/asish/linux subdir=tools/bpf/resolve_btfids --n  
o-print-directory -C bpf/resolve_btfids  
INSTALL libsubcmd_headers  
INSTALL libsubcmd_headers  
CALL scripts/checksyscalls.sh  
LDS scripts/module.lds  
CC init/main.o  
GEN scripts/gdb/linux/constants.py  
UPD init/utsversion-tmp.h  
CC init/do_mounts.o  
CC init/do_mounts_initrd.o  
CC init/initramfs.o  
CC init/calibrate.o  
CC init/initramfs.o
```

11. Now we reboot the VM for the new kernel to take place

```
Last login: Sun Dec 10 08:08:25 2023 from 68.250.225.193  
asish@cmpe283-vm3:~$ uname -a  
Linux cmpe283-vm3 6.7.0-rc4+ #3 SMP PREEMPT_DYNAMIC Sun Dec 10 09:19:41 UTC 2023 x86_64 x86_64 x86_64 GNU/Linux
```

12. Now since our kernel is in place we install required tools to initiate a innver VM running on a GCP-VM.

```

sudo apt install qemu-kvm
sudo apt install libvirt-clients
sudo apt install libvirt-daemon-system
sudo apt install bridge-utils
sudo apt install virt-manager

```

13. Once the required tools are installed, we now modify the code to report back some additional information with the CPUID leaf node. Change the cwd to `~/linux/arch/x86/kvm/vmx/vmx.c` and modify the file to add counters

```

extern atomic_t total_exits_counter;
extern atomic64_t total_cup_cycles_counter;

static int vmx_handle_exit(struct kvm_vcpu *vcpu, fastpath_t exit_fastpath)
{
    uint64_t begin_time_stamp_counter, end_time_stamp_counter;
    int ret;

    arch_atomic_inc(&total_exits_counter);
    begin_time_stamp_counter = rdtsc();

    ret = __vmx_handle_exit(vcpu, exit_fastpath);

    end_time_stamp_counter = rdtsc();

    arch_atomic64_add((end_time_stamp_counter - begin_time_stamp_counter),
&total_cup_cycles_counter);

}

```

14. we then modify the `~/linux/arch/x86/kvm/cpuid.c`

```

```c

int kvm_emulate_cpuid(struct kvm_vcpu *vcpu)
{
 u32 eax, ebx, ecx, edx;

 if (cpuid_fault_enabled(vcpu) && !kvm_require_cpl(vcpu, 0))
 return 1;

 eax = kvm_rax_read(vcpu);
 ecx = kvm_rcx_read(vcpu);

 switch(eax) {
 case 0x4FFFFFFC:
 eax = arch_atomic_read(&total_exits_counter);
 break;
 }
}

```

```

 case 0x4FFFFFFD:
 ebx = (atomic64_read(&total_cup_cycles_counter) >> 32);
 ecx = (atomic64_read(&total_cup_cycles_counter) & 0xFFFFFFFF);

 break;

 default:
 kvm_cpuid(vcpu, &eax, &ebx, &ecx, &edx, false);
 }

 kvm_rax_write(vcpu, eax);
}

...

```

15. Once the files are modified we again build our kernel to try it in inner VM.

```

```bash
make -j8 modules && make -j8 modules_install && make install
```

```

16. now we unload kvm kernel and load the kernel we just built using the following commands

```

rmmod arch/x86/kvm/kvm-intel.ko
rmmod arch/x86/kvm/kvm.ko
insmod arch/x86/kvm/kvm.ko
insmod arch/x86/kvm/kvm-intel.ko

```

17. Enable display driver inside GCP VM to get access of virt-manager which would install a vm inside

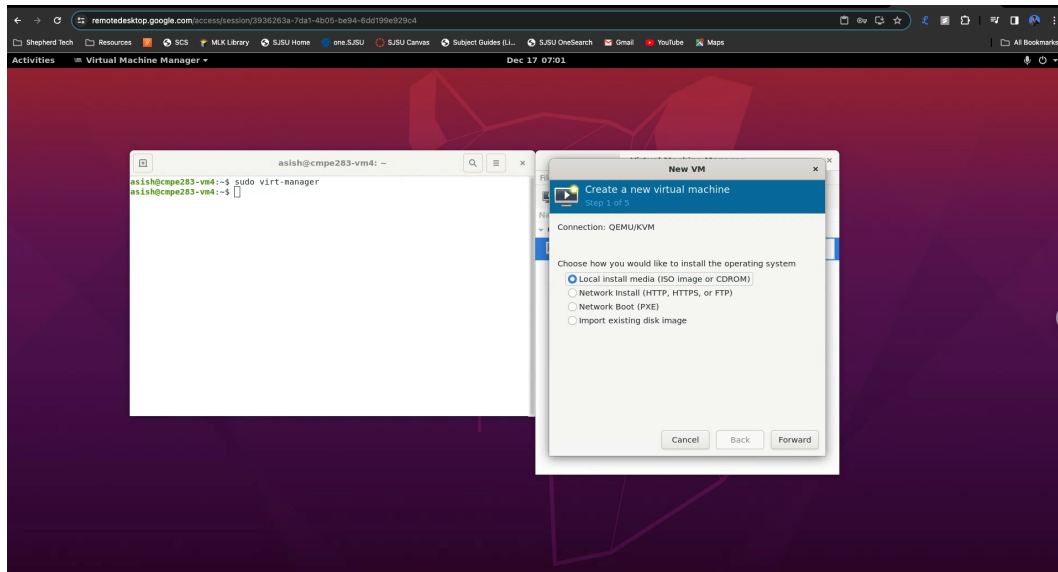
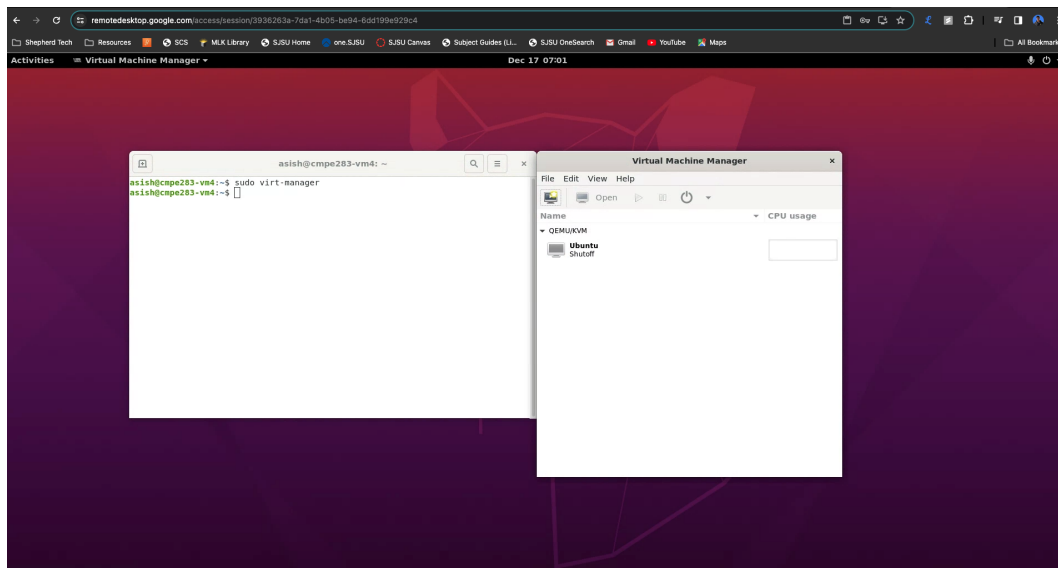
### Display device

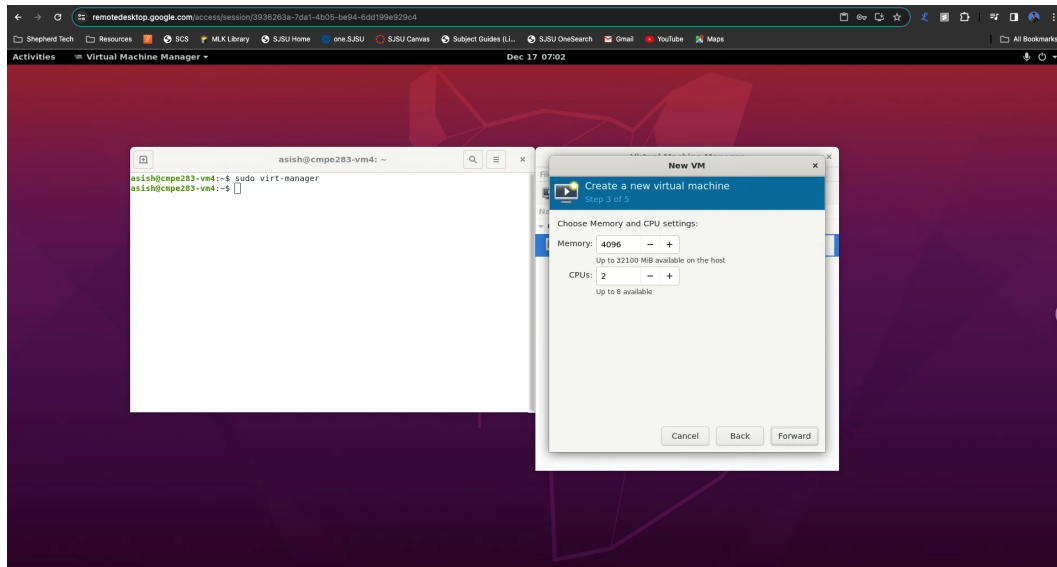
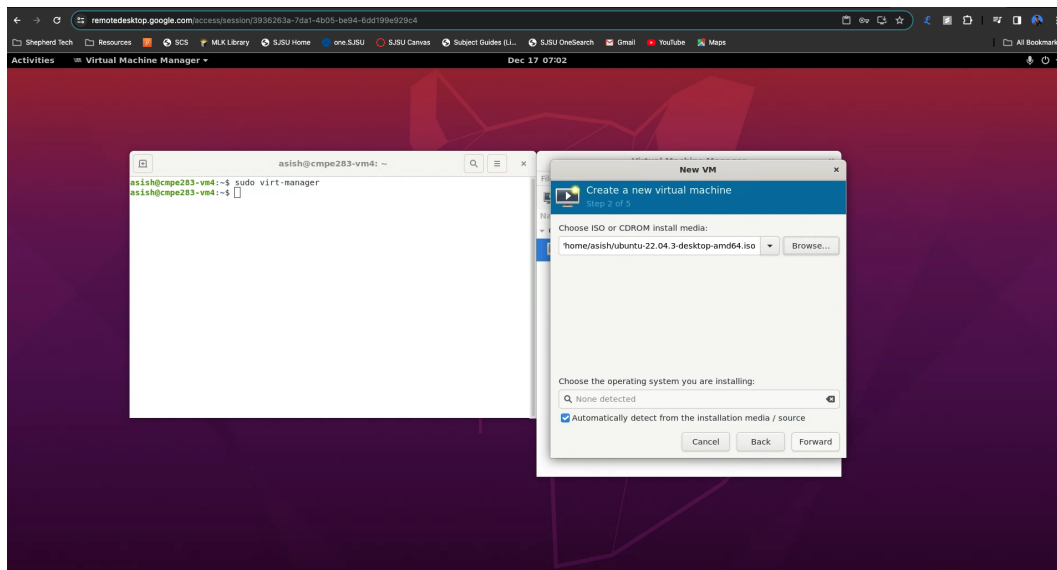
You must stop the VM instance to disable the display device.  
 Enable to use screen capturing and recording tools

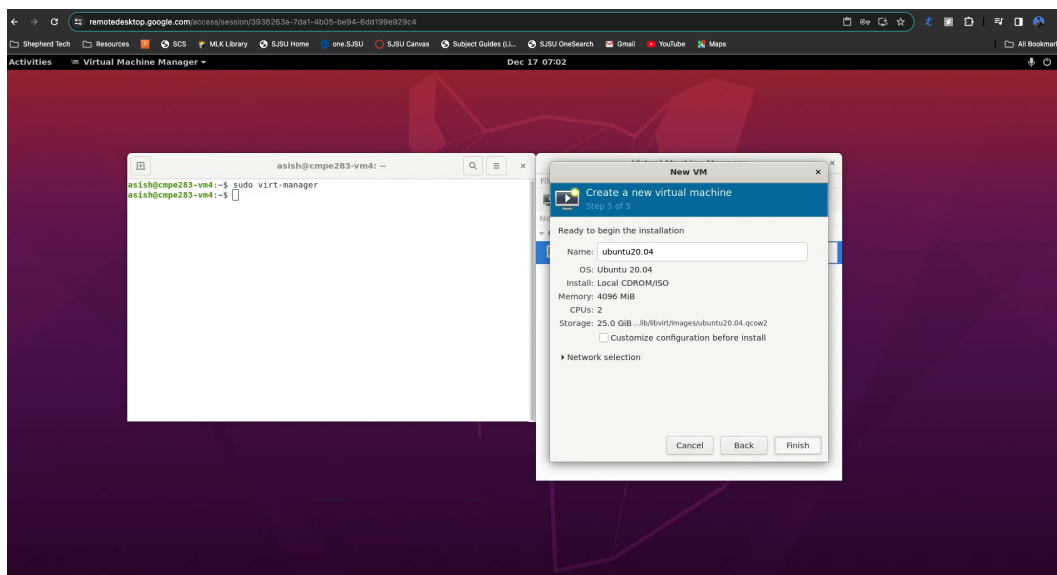
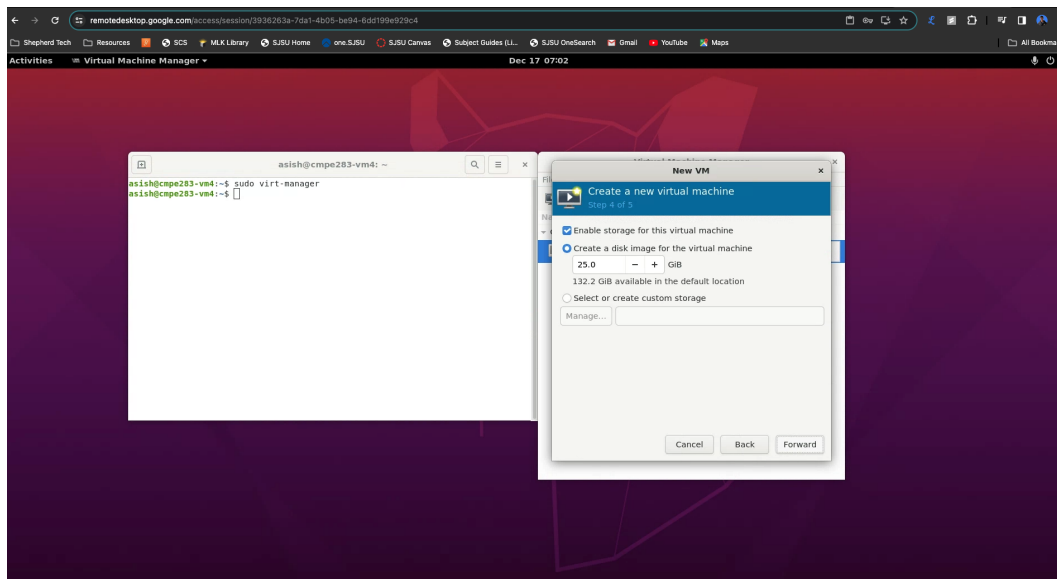
☒ Enable display device

it.

18. Using the chrom remote desktop console, we connect to our vm ip address to get display output and try to run `virt-manager` to install ubuntu.







19. Once the VM is setup we install ubuntu and install necessary tools

```
sudo apt-get update
sudo apt-get install git
sudo apt-get install make
sudo apt-get install build-essential
sudo apt-get install cpuid
```

20. Now we write a test .c file to test and its corresponding makefile and run the following command.

```
make
make run
```

21. The test output is as follows





Q1. Does the number of exits increase at a stable rate?

No, no of exits dont increase or decrease at a stable rate, some of stay the same, the delta between the increase and decrease is also not the same.

>> other questions answered in assignment 3